

**Exercise Sheet 8: Self Supervised Learning**

Due on 11.07.2025, 10:00

Nick Stracke (nick.stracke@lmu.de)

**Task 1: Temperature in Self Supervised Learning (3P)**

**Goal:** In this exercise, you will develop an intuitive understanding of how the *softmax* function works, and how the *temperature* parameter influences its behavior. The softmax function is widely used in machine learning to convert raw scores (e.g., similarities or logits) into probabilities. In self-supervised learning, temperature plays a key role in shaping how confident the model is in its predictions.

**1.1 Visualizing the Effect of Temperature**

Let us start with a simple vector of similarity scores:

$$\text{similarities} = [2.0, 1.0, 0.2, -0.5]$$

Implement the temperature-scaled softmax function:

$$\text{softmax}_{\tau}(x_i) = \frac{\exp(x_i/\tau)}{\sum_j \exp(x_j/\tau)}$$

- Plot the output of the softmax function for different temperature values:  $\tau \in \{0.1, 0.5, 1.0, 2.0, 5.0\}$ .
- Use a bar plot to visualize the resulting probability distribution for each  $\tau$ .

**1.2 Intuition: Confidence vs. Uncertainty**

Think of the softmax output as a way to distribute *attention* or *confidence* over a set of candidates (e.g., nearest neighbors in feature space). Answer the following:

- What happens to the distribution when the temperature is very low (e.g.,  $\tau = 0.1$ )? What does this mean in terms of model confidence?
- What happens when  $\tau$  is very high (e.g.,  $\tau = 10.0$ )? What does this mean for the model's ability to focus?

### 1.3 Why Temperature Matters in Contrastive Learning

In contrastive learning, the model must identify the correct (positive) example among many negative ones using similarity scores. Consider the following example:

`similarities = [1.0 (positive), 0.2, 0.1, -0.1, -0.3]`

- Compute the softmax probabilities for  $\tau = 10.0$  and  $\tau = 0.1$ .
- Which temperature gives the positive example the highest probability?
- Discuss: Why might a sharper distribution (lower  $\tau$ ) help the model learn more effectively?

### Task 2: SimCLR (14P)

**Goal:** In this exercise, you will implement a simplified version of SimCLR<sup>1</sup> for contrastive self-supervised learning. You will train an encoder to distinguish between differently augmented views of the same image and then evaluate the learned representations using a linear classifier and a t-SNE visualization.

#### 2.1 Implement and Train a SimCLR Model

- Take a look at `ex2.py`, which already includes a large part of the model and the training loop. This is your starting point.
- Finish the `ToDo's` in `SimCLRDataset`. Check the paper for what augmentations you need, but leave out Gaussian blur.
- Implement the contrastive loss. For a positive pair  $(i, j)$ , the loss is defined as

$$\ell_{i,j} = -\log \frac{\exp(\text{sim}(\mathbf{z}_i, \mathbf{z}_j)/\tau)}{\sum_{k=1}^{2N} \mathbb{1}_{[k \neq i]} \exp(\text{sim}(\mathbf{z}_i, \mathbf{z}_k)/\tau)}$$

with  $\mathbb{1}$  being the indicator function,  $\mathbf{z}$  the embedding and  $\tau$  the temperature. You can also have a look at Algorithm 1 in the paper for additional help.

- Train the model by calling the file directly. This is a standalone Python file because using `num_workers > 0` in a Jupyter notebook can cause issues. If you prefer working in a notebook, you might have to set `num_workers=0` if you encounter errors, but this often degrades performance. On an M1 Pro, training one epoch took roughly one minute when using mps.

<sup>1</sup><https://arxiv.org/abs/2002.05709>

## 2.2 Visualize with t-SNE

- We will use t-SNE to reduce the high-dimensional encoder features of the test set to 2D and visualize it.
- You have to finish the ToDo for extracting features.
- Don't worry if the result is not perfect. This is expected after only training for 10 epochs.

## 2.3 Quantitative Evaluation using Linear Probing

- To get a better feeling of the usefulness of the learned representation, we will perform classification on the frozen encoder features using a linear probe.
- Similar to the previous exercise, first extract the features for the entire train set.
- Then, define the linear probe, the optimizer, and the classification loss. Train the linear probe while keeping everything else frozen.
- If you trained the SimCLR model for 10 epochs and the linear probe for 50 epochs, you should achieve ca. 60% classification accuracy. For better performance, you have to train the base model for more epochs.

## Task 3: CLIP for Zero-Shot Classification (3P)

**Goal:** In this exercise, you will explore the CLIP model, which uses contrastive learning to align images and natural language. You will implement a simple zero-shot classifier by comparing the cosine similarity between image and text embeddings.

### 3.1 Understanding how CLIP Works

CLIP (Contrastive Language–Image Pretraining) is trained similarly to SimCLR: it maximizes the similarity between paired inputs. However, instead of comparing two augmented views of an image, CLIP learns to align **images and text descriptions** in a joint embedding space.

- Image and text encoders are trained to produce embeddings such that matching pairs (image, caption) have high cosine similarity.
- This enables CLIP to perform classification by comparing an image embedding to the embeddings of class name prompts, without needing any training on the target task.

### 3.2 Implement a Zero-Shot Classifier with CLIP

- Load the pretrained CLIP model from Hugging Face.
- Define a small set of unlabeled test images (e.g. 5–10).
- For each image, write down a set of possible text descriptions (e.g., “a photo of a dog”, “an image of a cat”, ...).
- Compute cosine similarities between the image embedding and each candidate text embedding.
- Assign the label of the text with the highest similarity.

### 3.3 Discussion

- Try varying your text prompts. Does the classification change?
- What kind of errors does CLIP make? Are they semantically reasonable?
- How does this compare to traditional classification models?

---

*Important Note: Submit exactly one ZIP file via Moodle before the deadline. Please submit your solution as a Jupyter Notebook. The ZIP file should contain your notebook and the images/videos you are loading. Make sure that it runs on different operating systems and uses relative paths. Non-trivial sections of your code should be explained with short comments, and variables should have self-explanatory names. The notebook file should contain your written code, all figures, explanations, and answers to questions.*